

Modeling Temporary Impact $g(x)$ from Market Data-

I modeled temporary price impact $g_t(x)$ as the execution-price premium (vs. mid) paid to trade size x at time t . Using top-10 MBP-10 snapshots for three tokens across 21 days, I constructed observations of $g(x)$ and compared non-linear, state-dependent models against classical parametric curves. I explored more flexible approaches that account for the concavity of the impact curve induced by the depth ladder and the strong dependence on spread, depth, imbalance, volatility, and time-of-day. Out-of-sample tests (GroupKFold by day/ticker) revealed that tree ensembles (Gradient Boosting, Random Forest) provided the most accurate predictions ($R^2 \approx 0.77$).

A. Constructing $g_t(x)$ from the order book (h.py)

Goal. For each snapshot t , I computed the slippage a market buy of size x would have incurred by sweeping the visible ask ladder (analogous for sells).

1. **Book features.** For each snapshot I computed the mid m_t , the best-level spread, the total top-10 depth V_t (both sides), the best-level imbalance, a rolling intraday volatility proxy, and the hour-of-day.
2. **Simulated sweep.** For sizes $x \in \{100, 200, \dots, 2000\}$, I accumulated fills from L1 upward and computed the average execution price $p^*_t(x)$.
3. **Temporary impact.** I defined $g_t(x) = p^*_t(x) - m_t$.

Output. I wrote an “enhanced slippage” CSV for each ticker-day with columns:

`timestamp, size, slippage (=gt(x)), vol_ratio (=x/depth), spread, depth, imbalance, volatility, hour_of_day`

These files fed the modeling stage.

Performance. I vectorized the sweep in a Numba-compiled loop over levels and sizes, which allowed me to process tens of millions of snapshots efficiently and write compact per-snapshot/per-size observations.

Non-linear, state-dependent modeling of $g_t(x)$ (I.py)

How I designed the study-

- **Target.** I modeled $g_t(x) = p^*_t(x) - m_t$ in price units.
- Used the liquidity ratio $rt = x/V_t$.
- **Features.** I engineered $\{x/V^{0.5}, x/V, \log x, \text{spread}, \text{depth}, \text{imbalance}, \text{volatility}, \text{hour_of_day}\}$
- **Aggregation.** I averaged each ticker-day into **20 size buckets** and tagged each with a `file_id`. Concatenating all days/tickers yielded $\sim 1,280$ rows.
- **Validation.** I used **GroupKFold (5 folds) by file_id**, so each test fold was a completely unseen day/ticker and there was no leakage.

Models I compared-

- **Parametric, non-linear in size/liquidity**
 - Square-root: $g = \sigma \sqrt{x/V}$
 - Logarithmic: $g = a + b \log x$
 - Quadratic in size: $g = b_1 x + b_2 x^2$
 - Power-law (liquidity-scaled): $g = \alpha (x/V)^\beta$ (non-linear fit for α, β)
- **Linear w/ regularization on full state**
 - Ridge, ElasticNet
- **Non-parametric, state-aware**
 - Random Forest, Gradient Boosting, XGBoost, KNN, SVR

Equations for Random Forest, Gradient Boost-

$$\hat{g}_t(x) = \frac{1}{B} \sum_{b=1}^B \sum_m c_{b,m} \mathbf{1}\{\phi_t(x) \in R_{b,m}\} \quad (\text{RF})$$

$$\hat{g}_t(x) = F_0 + \sum_{k=1}^K \sum_m \eta \gamma_k d_{k,m} \mathbf{1}\{\phi_t(x) \in R_{k,m}\} \quad (\text{GB})$$

with $\phi_t(x) = [\sqrt{x/V_t}, x/V_t, \log x, \text{spread}_t, V_t, \text{imbalance}_t, \text{vol}_t, \text{hour}_t]$.

Results-

Cross-validated Model Performance:

model	mean_test_mse	std_test_mse	mean_test_r2	std_test_r2
GradientBoost	2.389680	0.835836	0.774102	0.074274
RandomForest	2.463870	1.038403	0.766079	0.100258
XGBoost	2.834322	0.827539	0.725054	0.099264
Ridge	4.420681	1.165765	0.589537	0.054985
Linear x/V	5.726270	1.539792	0.470070	0.070683
PowerLaw x/V	5.733184	1.579140	0.469843	0.075209
Square-root	6.728375	1.438590	0.373448	0.039820
Quadratic	6.944874	1.335428	0.349971	0.041279
ElasticNet	7.683910	1.540505	0.284034	0.026587
Logarithmic	8.325828	1.400881	0.218657	0.031837
SVR	10.848014	2.022484	-0.012478	0.035985

Interpretation and Model Choice-

Top predictive accuracy-

- **Gradient Boosting** achieved the **lowest MSE (2.39)** and **highest mean $R^2=0.774$** , with a moderate fold-to-fold variability ($\sigma R^2=0.074$).
- **Random Forest** was a close second ($R^2=0.766$, but higher variance), and **XGBoost** followed ($R^2=0.725$).

Gradient Boosting is the best out-of-sample predictor of $gt(x)$. It flexibly captures both the concave size-dependence and the interactions with spread, depth, imbalance, volatility, and time-of-day. To guard against overfitting, I used 5-fold GroupKFold cross-validation by day/ticker and kept tree depth and learning-rate settings moderate (e.g. 100 trees, default max depth), ensuring the model generalizes across unseen days rather than memorizing noise. Its CV variability is modest, giving confidence in its stability.